

Process Relationship

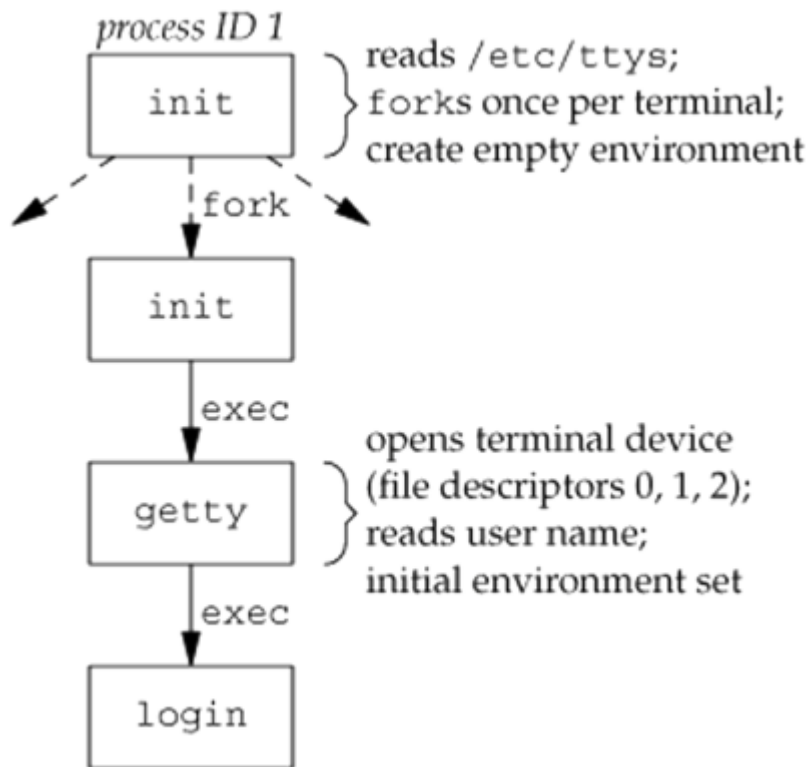
Table of Contents

- [Process Relationship](#)
 - [Table of Contents](#)
 - [Login](#)
 - [Process Groups](#)
 - [Sessions](#)
 - [Controlling Terminal](#)
 - [Job Control](#)
 - [Shell Execution of Programs](#)
 - [Orphaned Process Groups](#)

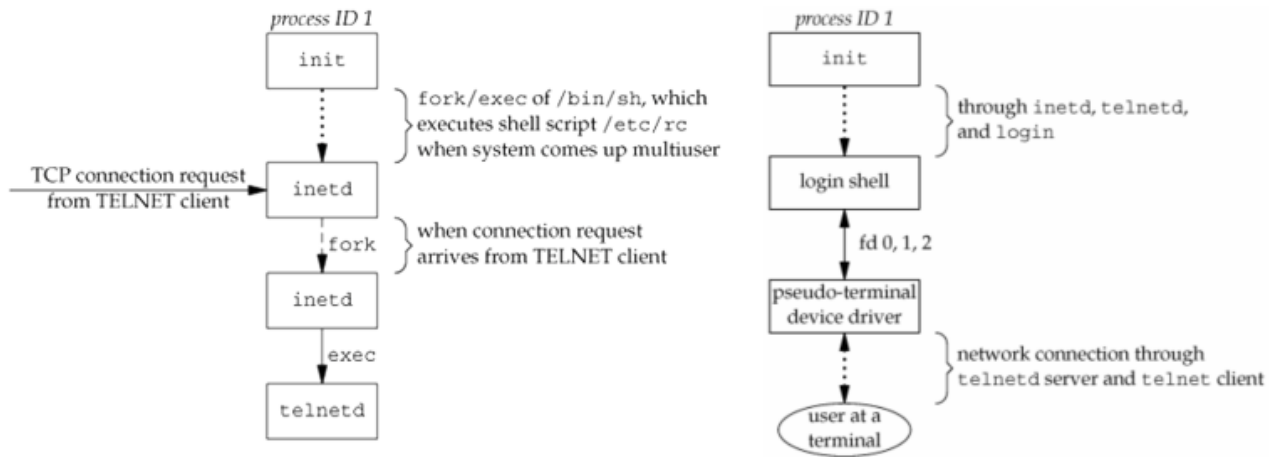
Login

- Linux Booting Process
 - first process: `/sbin/init`
 - PID is 1
 - configurations: `/etc/inittab`, `/etc/event.d/*`, or `/etc/init/*`
 - run levels
 - 0: halted
 - 6: reboot
 - 1 to 5 can be customized
 - default run level often set to 2, 3 or 5
 - enable console logs
- Linux Terminal Logins (e.g. Ubuntu)
 - start 6 consoles terminals for login
 - can be switched by Alt + F1 to F6
 - `getty` program
 - call open for `/dev/tty1`, `/dev/ttyS0`, ...
 - create file descriptor 0, 1 and 2
 - show login prompt
 - when user provide name, invoke `/bin/login` program
 - `execle("/bin/login", "login", "-p", username, (char *)0, envp);`
 - `login` program
 - display `Password` prompt
 - read password by `getpass(3)`
 - read encrypted password
 - encrypt input and compare password
 - if fails, the program terminates and `init` restart `getty`
 - after successful `login`

- `chdir` to user's home directory
- `chown` the terminal device owner to user
- set the access permissions for the terminal device
- `setgid` set group IDs
- `initgroups` initialize the supplementary group access list
- initialize the environment variables
- `setuid` set user ID and invoke a login shell



- Network Logins – via the `telnetd` Program
 - open a pseudo-terminal device
 - `/dev/pts/N`
 - splits into two processes using `fork`
 - parent handles the communication across the network connection
 - child does an `exec` of the `login` program
 - whether we log in through a terminal or a network connection ...
 - we have a login shell
 - `std in/out/err` are connected to either a terminal device or a pseudo-terminal device



Process Groups

- Definition
 - a collection of one or more processes, usually same jobs
- Properties
 - parent process can wait children in a process group
 - signals can be sent to processes in a process group
- `getpgid(2): pid_t getpgid(pid_t pid);`
- `getpgrp(2): pid_t getpgrp(void); = getpgid(0)`
- Group Leader
 - leader process group ID equals its process ID
 - leader can create a group, create processes in the group, and then quit
- Process Group Lifetime
 - start on the creation of the group
 - end when the last process in the group leaves
- Create/Join a Process Group
 - `setpgid(2): int setpgid(pid_t pid, pid_t pgid);`
 - set process group ID of process `pid` to `pgid`
 - if `pid` equals `pgid`, then process `pid` becomes leader
 - if `pid` equals 0, then `pid` is replaced by the caller `pid`
 - if `pgid` equals 0, then `pgid` is replaced by the value `pid`
 - `setpgid` limitations
 - process can only perform `setpgid` to itself or its children
 - process can not perform `setpgid` after its child calls `exec`
 - use of `setpgid` function
 - called after `fork`, parent sets child process group ID
 - child sets its own process group ID
 - above two actions are redundant, but they guaranteed that the child is placed into its own process group

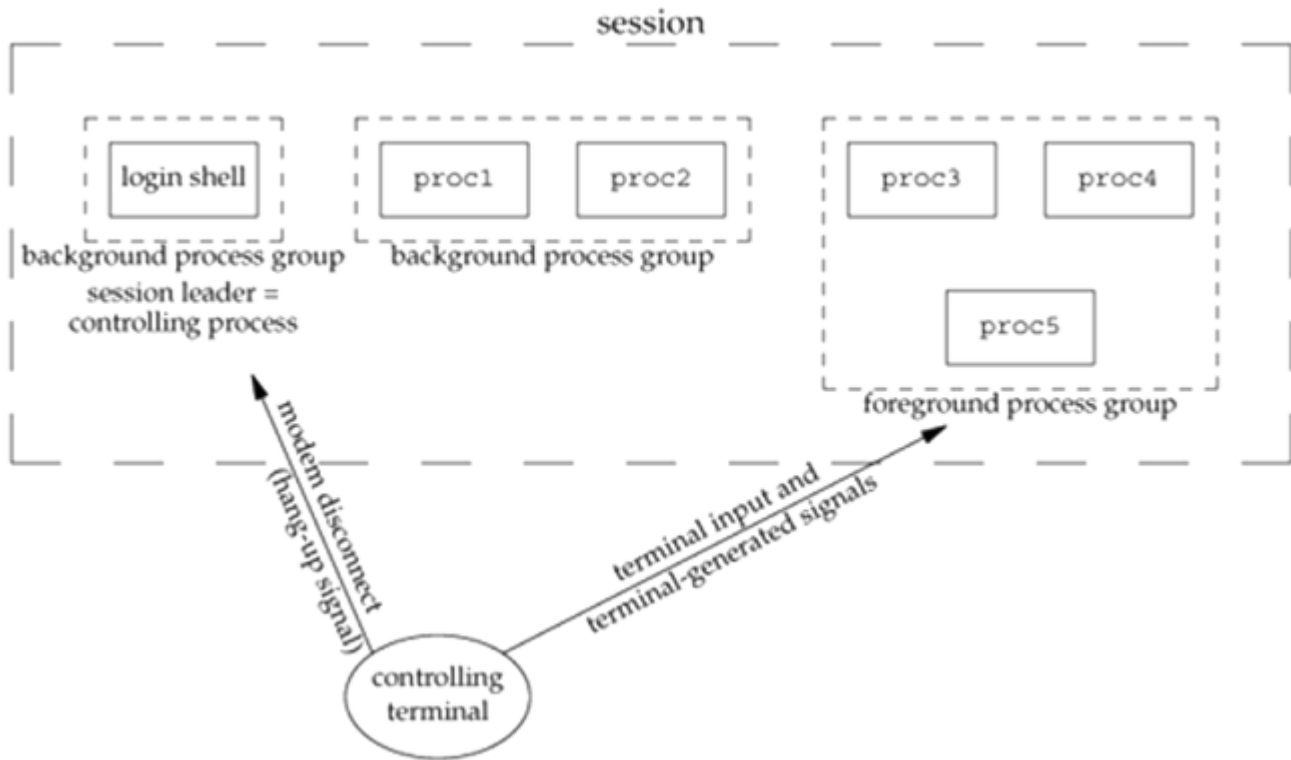
Sessions

- Definition
 - a collection of one or more process groups
 - pipes are used to generate process group

- `$ prog1 | prog2 | prog3`
- these three programs are in the same process group
- Create a Session
 - `setsid(2): pid_t setsid(void);`
 - return: `pgid` or -1 if caller is process group leader
 - if the process is not a process group leader, the function creates a new session
 - the process becomes session leader
 - the process is the only process in the session
 - the process becomes the process group leader of a new process group
 - the new process group ID is the process ID of the calling process
 - the process has no controlling terminal
- Get Current Session ID
 - `getsid(2): pid_t getsid(pid_t pid);`
 - return: the session leader's process group ID, -1 on error
 - if `pid` is 0, `pid` is replaced by caller process ID
- when user login, the session leader is usually the shell

Controlling Terminal

- a session can have a single controlling terminal
 - terminal device or pseudo-terminal device
 - the session leader that establishes the connection to the controlling terminal is called the controlling process
- the process groups within a session can be divided into:
 - a single foreground process group
 - one or more background process groups
- if a session has a controlling terminal
 - it has a single foreground process group
 - all other process groups in the session are background process groups
- User control keys
 - send signals to all processes in the foreground process group
 - interrupt key (often Ctrl + C): send `SIGINT`
 - quit key (often Ctrl + Backspace): send `SIGQUIT`
 - if a network disconnect is detected by the terminal interface, the `SIGHUP` is sent to the controlling process



- Whom to Send Signals
 - `tcgetpgrp(3): pid_t tcgetpgrp(int fd);`
 - get terminal foreground process group
 - `tcsetpgrp(3): int tcsetpgrp(int fd, pid_t pgrp);`
 - only set by controlling process
 - most applications don't call above two functions
 - normally called by job-control shells
- Direct Access to the Controlling Terminal
 - even stdin or stdout is redirected
 - this can be done by open file `/dev/tty`
 - synonym within the kernel for the controlling terminal
 - the program must have a controlling terminal, otherwise open will fail

Job Control

- start multiple jobs from a single terminal
- requirements
 - a shell that supports job control
 - the terminal driver in the kernel must support job control
 - the kernel must support certain job-control signals
- start a job in background: `&` operator
- stop job running in foreground
 - Ctrl + z
 - `SIGTSTP` is sent to all processes in the foreground process group
- `SIGTTIN` and `SIGTTOU`
 - processes in the foreground process group is always able to read from and write to the terminal
 - but background process is restricted to do so
 - signal will send to the background process

- **SIGTTIN**: signal of reading from terminal

```
$ cat > temp.foo &           start in background, but it'll read from standard input
[1] 1681
$                               we press RETURN
[1] + Stopped cat > temp.foo
$ fg %1                       bring job number 1 into the foreground
cat > temp.foo                the shell tells us which job is now in the foreground
hello, world                  enter one line
^D                             type the end-of-file character
$ cat temp.foo                check that the one line was put into the file
hello, world
```

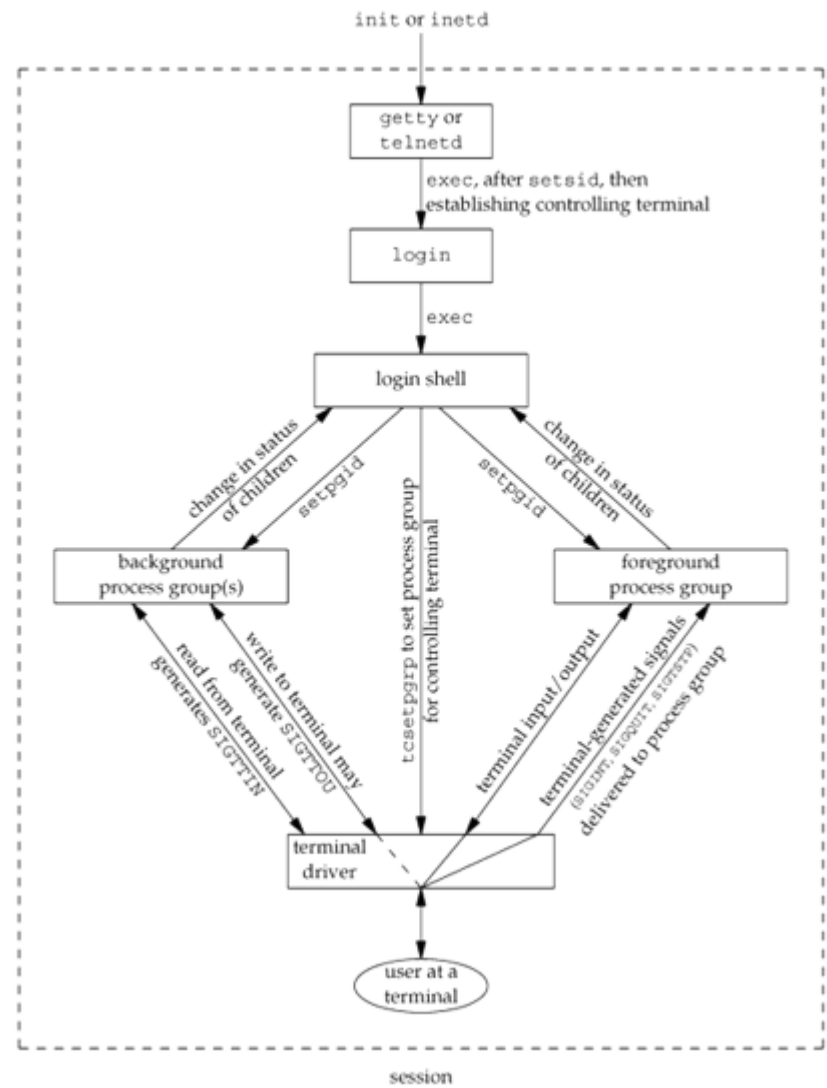
- **SIGTTOU**: signal of writing to terminal

```
$ cat temp.foo &              execute in background
[1] 1719
$ hello, world                the output from the background job appears after the prompt
                               we press RETURN

[1] + Done cat temp.foo
$ stty tostop                  disable ability of background jobs to output
                               to the controlling terminal

$ cat temp.foo &              try it again in the background
[1] 1721
$                               we press RETURN and find the job is stopped
[1] + Stopped(SIGTTOU) cat temp.foo
$ fg %1                       resume stopped job in the foreground
cat temp.foo                  the shell tells us which job is now in the foreground
hello, world                  and here is its output
```

- job control summary



Shell Execution of Programs

- `ps -o pid,ppid,pgid,sid,comm | cat1 | cat2`

```
PID  PPID  PGID  SID  COMMAND
949   947   949   949   sh
1988  949   949   949   cat2
1989  1988  949   949   ps
1990  1988  949   949   cat1
```

Orphaned Process Groups

- Definition
 - the parent process of every member is either
 1. a member of the group or
 2. not a member of the group's session
- In contrast, process group is not orphaned if

- a process in the group has a parent in a different process group but in the same session
- If a process group becomes orphaned
 - every stopped process in the group is sent the **SIGHUP** followed by the **SIGCONT**
 - the default action on receipt of a **SIGHUP** is to terminate the process
- Example

```
int main() {
    char c;
    pid_t pid;
    pr_ids("parent");
    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid > 0) {
        sleep(5);
    } else {
        pr_ids("child");
        signal(SIGHUP, sig_hup); /* establish signal handler */
        kill(getpid(), SIGTSTP); /* stop ourself */
        pr_ids("child");          /* prints only if we're continued */
        if (read(STDIN_FILENO, &c, 1) != 1)
            printf("read error %d on controlling TTY\n", errno);
    }
    exit(0);
}
```

```
parent: pid = 6099, ppid = 2837, pgrp = 6099, tpgrp = 6099
child: pid = 6100, ppid = 6099, pgrp = 6099, tpgrp = 6099
(sleep for 5 seconds)
SIGHUP received, pid = 6100
child: pid = 6100, ppid = 1, pgrp = 6099, tpgrp = 2837
read error from controlling TTY, errno = 5
```

1. the parent and the child prints out their own information
2. the parent then sleeps for 5 seconds
3. the child stopped itself
4. when the parent terminates, the child received **SIGHUP** and **SIGCONT**
5. since the child has assigned the **SIGHUP** handler, it is not terminated
6. the child is now in background, so read from TTY got the **EIO** error