

# Signals

---

## Table of Contents

- [Signals](#)
  - [Table of Contents](#)
  - [Introduction](#)
  - [The Signal Function](#)
  - [Interrupted System Calls](#)
  - [Reentrant Functions](#)
  - [Reliable Signal Model](#)
  - [Functions for Handling Signals](#)
  - [setjmp\(3\) and longjmp\(3\)](#)
  - [The System Function](#)

## Introduction

- Signals are Software Interrupts to Handle Asynchronous Events
- [signal\(7\)](#)
- Signal Concepts
  - numerous conditions
    - the terminal-generated signals
    - hardware exceptions
    - software conditions
    - the [kill\(2\)](#) function
    - the [kill\(1\)](#) command
  - handle signal
    - ignore the signal
      - but we cannot ignore [SIGKILL](#) and [SIGSTOP](#)
    - catch the signal, the register customized handler
      - but we cannot register [SIGKILL](#) and [SIGSTOP](#)
    - let the default action apply
      - every signal has a default action
- List of Linux Signals (e.g. Ubuntu 18.04 x86\_64)
  - POSIX.1-1990 Standard

Signal	Value	Action	Comment
<a href="#">SIGHUP</a>	1	Term	Hangup detected on controlling terminal or death of controlling process
<a href="#">SIGINT</a>	2	Term	Interrupt from keyboard
<a href="#">SIGQUIT</a>	3	Core	Quit from keyboard
<a href="#">SIGILL</a>	4	Core	Illegal Instruction
<a href="#">SIGABRT</a>	6	Core	Abort signal from <a href="#">abort(3)</a>

Signal	Value	Action	Comment
<b>SIGFPE</b>	8	Core	Floating-point exception
<b>SIGKILL</b>	9	Term	Kill signal
<b>SIGSEGV</b>	11	Core	Invalid memory reference
<b>SIGPIPE</b>	13	Term	Broken pipe: write to pipe with no readers; see <b>pipe(7)</b>
<b>SIGALRM</b>	14	Term	Timer signal from <b>alarm(2)</b>
<b>SIGTERM</b>	15	Term	Termination signal
<b>SIGUSR1</b>	30,10,16	Term	User-defined signal 1
<b>SIGUSR2</b>	31,12,17	Term	User-defined signal 2
<b>SIGCHLD</b>	20,17,18	Ign	Child stopped or terminated
<b>SIGCONT</b>	19,18,25	Cont	Continue if stopped
<b>SIGSTOP</b>	17,19,23	Stop	Stop process
<b>SIGTSTP</b>	18,20,24	Stop	Stop typed at terminal
<b>SIGTTIN</b>	21,21,26	Stop	Terminal input for background process
<b>SIGTTOU</b>	22,22,27	Stop	Terminal output for background process

- SUSv2 and POSIX.1-2001

Signal	Value	Action	Comment
<b>SIGBUS</b>	10,7,10	Core	Bus error (bad memory access)
<b>SIGPOLL</b>		Term	Pollable event (Sys V). Synonym for <b>SIGIO</b>
<b>SIGPROF</b>	27,27,29	Term	Profiling timer expired
<b>SIGSYS</b>	12,31,12	Core	Bad system call (SVr4); see also <b>seccomp(2)</b>
<b>SIGTRAP</b>	5	Core	Trace/breakpoint trap
<b>SIGURG</b>	16,23,21	Ign	Urgent condition on socket (4.2BSD)
<b>SIGVTALRM</b>	26,26,28	Term	Virtual alarm clock (4.2BSD)
<b>SIGXCPU</b>	24,24,30	Core	CPU time limit exceeded (4.2BSD); see <b>setrlimit(2)</b>
<b>SIGXFSZ</b>	25,25,31	Core	File size limit exceeded (4.2BSD); see <b>setrlimit(2)</b>

- Linux 2.4

Signal	Value	Action	Comment
<b>SIGIOT</b>	6	Core	IOT trap. A synonym for <b>SIGABRT</b>
<b>SIGEMT</b>	7,-,7	Term	Emulator trap
<b>SIGSTKFLT</b>	-,16,-	Term	Stack fault on coprocessor (unused)

Signal	Value	Action	Comment
<b>SIGIO</b>	23,29,22	Term	I/O now possible (4.2BSD)
<b>SIGCLD</b>	-, -,18	Ign	A synonym for <b>SIGCHLD</b>
<b>SIGPWR</b>	29,30,19	Term	Power failure (System V)
<b>SIGINFO</b>	29,-,-		A synonym for <b>SIGPWR</b>
<b>SIGLOST</b>	-, -, -	Term	File lock lost (unused)
<b>SIGWINCH</b>	28,28,20	Ign	Window resize signal (4.3BSD, Sun)
<b>SIGUNUSED</b>	-,31,-	Core	Synonymous with <b>SIGSYS</b>

## The Signal Function

- `signal(2): void (*signal(int signum, void (*handler)(int)))(int)`
  - `typedef void (*sighandler_t)(int);`
  - `sighandler_t signal(int signum, sighandler_t handler);`
  - return: previous disposition of signal, **SIGERR** on error
  - `signum` is the name of the signal
  - `handler` is the function to be called when the signal occurs
    - can also be **SIG\_IGN** or **SIG\_DFL**
    - **SIG\_IGN** ignore `signum`
    - **SIG\_DFL** use default action in above table
  - implementations differs among system
    - better to use `sigaction` function
- Example

```
void sig_usr(int signo) {
    if (signo == SIGUSR1) {
        printf("received SIGUSR1\n");
    } else if (signo == SIGUSR2) {
        printf("received SIGUSR2\n");
    } else {
        printf("receive signal %d\n", signo);
    }
}

int main() {
    if (signal(SIGUSR1, sig_usr) == SIG_ERR) {
        printf("can't catch SIGUSR1");
    }
    if (signal(SIGUSR2, sig_usr) == SIG_ERR) {
        printf("can't catch SIGUSR2");
    }
    while (1) {
        pause();
    }
}
```

```
}
}
```

```
$ ./a.out &
[1] 15512
$ kill -USR1 15512
received SIGUSR1
$ kill -USR2 15512
received SIGUSR2
$ kill 15512 # default send SIGTERM
[1]+ Terminated ./a.out
$ kill -l # check available signals
1) SIGHUP 2) SIGINT ...
```

- Signal Setup
  - signal is either default or ignore when program is executed
  - `fork(2)`: child inherits parent's signal dispositions
  - `exec(3)`: child changes the signal dispositions to their default action
  - the shell sets `SIGINT` and `SIGQUIT` in the background process to be ignored

## Interrupted System Calls

- A System Call may be Interrupted
  - if process catch a signal while in slow system call
  - then the system call return an error and `errno` was set to `EINTR`
    - `EINTR`: 4, interrupted system call
- Handle Interrupted System Calls

```
again:
if ((n = read(fd, buf, BUFSIZE)) < 0) {
    if (errno == EINTR)
        goto again;
}
```

- Prevent Interrupted System Calls
  - 4.2BSD introduced automatic restarting of certain interrupted system calls
  - `ioctl`, `read`, `readv`, `write`, `writv`, `wait`, and `waitpid`
  - the feature can be disabled if you don't like it

## Reentrant Functions

- When a Signal is Being Caught

- instruction is interrupted
- execute signal handler instructions
- if the handler returns, then the process continues the interrupted instructions
- If Handler Uses the Same System Call as the Interrupted Instructions
  - if works well, they are reentrant functions
  - some functions can not work with reentrant, because
    - they use static data structures
    - they are part of the standard I/O library (buffering)
    - they call `malloc` or `free`
  - reentrant functions example

```
void handler(int signo) {
    tmpnam(NULL);
}
int main() {
    char *s = tmpnam(NULL);
    signal(SIGALRM, handler);
    alarm(3);
    for (int i = 0; i < 6; ++i) {
        printf("tmpnam = %s\n", s);
        sleep(1);
    }
}
```

- `SIGCLD` and `SIGCHLD`
  - the signal disposition is `SIG_DFL`: default ignore
  - zombie avoidance
    - explicitly set disposition to `SIG_IGN`
    - no zombie will be created if children of the calling process terminate

## Reliable Signal Model

- Delivery of a Signal
  - delivered: a process received the signal and action is taken
  - pending: a signal is generated, but not delivered
- Blocking the Delivery of a Signal
  - if signal is blocked and handler is `SIG_DFL` or a handler
  - then the signal remains pending until
    - unblock the signal or
    - change handler to `SIG_IGN`
- Block Signal is Generated More Than Once
  - POSIX.1 allows deliver signal once or more
  - most UNIX systems do not queue signals

- Deliver Signals More Than Once
  - POSIX.1 does not specify the order
  - it suggests that signals related to the current state be delivered before other signals
- Signal Mask
  - defines the blocked signals in a process
  - each signal has a corresponding bit

## Functions for Handling Signals

- `kill(2): int kill(pid_t pid, int sig);`
- `raise(3): int raise(int sig);`
  - return: 0 OK, -1 error
  - send signal to
    - `pid > 0`: pid process
    - `pid == 0`: every process in the process group as same as the caller
    - `pid == -1`: every process the caller has permission to send, except process 1
    - `pid < -1`: every process in the process group `-pid`
  - `raise(sig)` equals `kill(getpid(), sig)`
  - NULL signal
    - may be used to check the existence of a process
    - no such process: `kill` returns -1 and `errno` is set to `ESRCH`
      - `ESRCH`: 3, no such process
- `alarm(2): unsigned int alarm(unsigned int seconds);`
  - return: 0 or number of seconds remaining until previously scheduled alarm
  - `SIGALRM` is generated when the timer expires, default terminate the process
  - only one alarm clocks per process
- `pause(2): int pause(void);`
  - return: -1 with `errno` set to `EINTR`
  - suspend a process until it received a signal
- `abort(3): void abort(void);`
  - send `SIGABRT` to the caller
  - the signal can be caught by a handler, but
    - the handler will not return if it calls `exit(3)`, `_exit(2)` or `_Exit(2)`
    - if the handler returns, then `abort` terminates the process
    - if the handler calls `siglongjmp`, the program may continue to execute
- Signal Sets: `sigsetops(3)`
  - operation
    - `sigemptyset(3): int sigemptyset(sigset_t *set);`
      - initializes `set` to empty
    - `sigfillset(3): int sigfillset(sigset_t *set);`

- initializes `set` to full
  - `sigaddset(3): int sigaddset(sigset_t *set, int signum);`
    - add `signum` to `set`
  - `sigdelset(3): int sigdelset(sigset_t *set, int signum);`
    - delete `signum` from `set`
  - return: 0 OK, -1 error
- membership test
  - `sigismember(3): int sigismember(const sigset_t *set, int signum);`
  - return: 1 if true, 0 if false, -1 error
- `sigprocmask(2): int sigprocmask(int how, const sigset_t *set, sigset_t *oldset);`
  - return: 0 OK, -1 error
  - block or unblock signals
  - `how`
    - `SIG_BLOCK`: `$newblkset = curblkset \cup set$`
    - `SIG_UNBLOCK`: `$newblkset = curblkset - set$`
    - `SIG_SETMASK`: `$newblkset = set$`
  - if `oldset` is not `NULL`, then previous value of the signal mask is stored `oldset`
  - if `set` is `NULL`, then the mask remains unchanged
- `sigpending(2): int sigpending(sigset_t *set);`
  - return: 0 OK, -1 error
  - get currently pending signals
- Example

```
int main() {
    sigset_t newmask, oldmask, pendmask;
    /* sig_quit print caught SIGQUIT and then reset SIGQUIT handler */
    if (signal(SIGQUIT, sig_quit) == SIG_ERR) {
        printf("err signal register\n");
    }
    sigemptyset(&newmask);
    sigaddset(&newmask, SIGQUIT);
    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) {
        printf("err sigprocmask\n");
    }
    printf("SIGQUIT blocked\n");
    sleep(5);
    if (sigpending(&pendmask) < 0) {
        printf("err sigpending\n");
    }
    if (sigismember(&pendmask, SIGQUIT)) {
        printf("SIGQUIT in pending\n");
    }
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) {
        printf("err sigprocmask\n");
    }
}
```

```

    }
    printf("SIGQUIT unblocked\n");
    sleep(5);
}

```

```

$ ./a.out
^\\^\\^\\
SIGQUIT in pending
caught SIGQUIT
SIGQUIT unblocked

```

- `sigaction(2):int sigaction(int signum, const struct sigaction *act, struct sigaction *oldact);`
  - return: 0 OK, -1 error
  - examine or modify (or both) the signal action
  - `sigaction`

```

struct sigaction {
    void (*sa_handler)(int); // default handler
    void (*sa_sigaction)(int, siginfo_t *, void *); // alternative
    handler when SA_SIGINFO is enabled
    sigset_t sa_mask; // blocked signals when executing the
    handler
    int sa_flags; // various options
    void (*sa_restorer)(void);
};

```

- `sa_flags`

<code>sa_flags</code>	Description
<code>SA_INTERRUPT</code>	obsoleted after Linux 2.6.19
<code>SA_NOCLDSTOP</code>	when <code>signum</code> is <code>SIGCHLD</code> , do not receive notification when child processes stop
<code>SA_NOCLDWAIT</code>	when <code>signum</code> is <code>SIGCHLD</code> , do not transform children into zombies
<code>SA_RESETHAND</code>	when handler is established, restore the action to the default upon entry to the handler
<code>SA_RESTART</code>	when handler is established, restart some system calls instead of generating <code>EINTR</code>
<code>SA_NODEFER</code>	when handler is established, do not prevent the signal from being received from within its own handler



sa_flags	Description
SA_SIGINFO	when handler is established, use <code>sa_sigaction</code> handler
SA_ONSTACK	when handler is established, call the handler on an alternate signal stack provided by <code>sigaltstack(2)</code>
SA_RESTORER	used in C libraries, <code>sa_restorer</code> field contains the address of a "signal trampoline", see <code>sigreturn(2)</code>

- `sigsuspend(2): int sigsuspend(const sigset_t *mask);`
  - return: -1 with `errno` set to `EINTR`
  - replace signal mask and then suspends the process until
    - a signal whose action is to invoke a handler or
    - a signal to terminate a process, `sigsuspend(2)` does not return in this condition
  - it is not possible to block `SIGKILL` or `SIGSTOP`
  - catching signal between `sigprocmask(2)` and `pause(2)`
    - after handler is done, the program is suspended
    - use `sigsuspend(2)` to prevent this condition
- Example: Wait for a Global Variable to be Set

```
volatile sig_atomic_t quitflag;

void sig_int(int signo) {
    if (signo == SIGINT)
        printf("\ninterrupt\n");
    else if (signo == SIGQUIT)
        quitflag = 1;
}

int main(void) {
    sigset_t newmask, oldmask, zeromask;
    if (signal(SIGINT, sig_int) == SIG_ERR) printf("signal(SIGINT)
error");
    if (signal(SIGQUIT, sig_int) == SIG_ERR) printf("signal(SIGQUIT)
error");
    sigemptyset(&zeromask);
    sigemptyset(&newmask);
    sigaddset(&newmask, SIGQUIT);
    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
printf("SIG_BLOCK error");
    while (quitflag == 0) sigsuspend(&zeromask);
    quitflag = 0;
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
printf("SIG_SETMASK error");
}
```

- Example: Process Synchronization

```

// library
volatile sig_atomic_t sigflag;
sigset_t newmask, oldmask, zeromask;
void sig_usr(int signo) { sigflag = 1; }

void TELL_WAIT(void) {
    if (signal(SIGUSR1, sig_usr) == SIG_ERR) printf("signal(SIGUSR1)
error");
    if (signal(SIGUSR2, sig_usr) == SIG_ERR) printf("signal(SIGUSR2)
error");
    sigemptyset(&zeromask);
    sigemptyset(&newmask);
    sigaddset(&newmask, SIGUSR1);
    sigaddset(&newmask, SIGUSR2);
    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
printf("SIG_BLOCK error");
}

void TELL_PARENT(pid_t pid) { kill(pid, SIGUSR2); }

void TELL_CHILD(pid_t pid) { kill(pid, SIGUSR1); }

void WAIT_PARENT(void) {
    while (sigflag == 0) sigsuspend(&zeromask);
    sigflag = 0;
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
printf("SIG_SETMASK error");
}

void WAIT_CHILD(void) {
    while (sigflag == 0) sigsuspend(&zeromask);
    sigflag = 0;
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
printf("SIG_SETMASK error");
}

// application
int main(void) {
    pid_t pid;
    TELL_WAIT();
    if ((pid = fork()) < 0) {
        printf("fork error");
    } else if (pid == 0) {
        WAIT_PARENT(); /* parent goes first */
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
        TELL_CHILD(pid);
    }
}

```

## setjmp(3) and longjmp(3)

- `setjmp(3): int setjmp(jmp_buf env);`
- `longjmp(3): void longjmp(jmp_buf env, int val);`
- `setjmp(3)` and `longjmp(3)` can Jump Across Functions
  - handler perform `longjmp(3)` to main loop
  - but mask of a signal is set automatically when that signal is caught
  - `longjmp(3)` exits the handler with the mask, and it will cause problems
- Solutions
  - `setjmp(3)` should save signal mask
  - the mask is restored when `longjmp(3)` is called
- POSIX.1 does not define how do `setjmp(3)` and `longjmp(3)` handle masks
  - FreeBSD and MAC OS X save and restore the mask automatically
  - however, Linux must work with `sigsetjmp(3)` and `siglongjmp(3)`
- `sigsetjmp(3): int sigsetjmp(sigjmp_buf env, int savesigs);`
  - return: 0 if called directly, nonzero if returning from a call to `siglongjmp(3)`
  - if `savesigs` is nonzero
    - usually set `savesig` to nonzero
    - the process's current signal mask is saved in `env`
    - the mask will be restored if a `siglongjmp(3)` is later performed with this `env`
- `siglongjmp(3): void siglongjmp(sigjmp_buf env, int val);`
- Example

```
sigjmp_buf jmpbuf;
volatile sig_atomic_t canjump;

void sig_usr1(int signo) {
    time_t starttime;
    if (canjump == 0) return;
    pr_mask("starting sig_usr1: ");
    alarm(3);
    starttime = time(NULL);
    for (;;)
        if (time(NULL) > starttime + 5) break;
    pr_mask("finishing sig_usr1: ");
    canjump = 0;
    siglongjmp(jmpbuf, 1);
}

void sig_alrm(int signo) { pr_mask("in sig_alrm: "); }
```

```

int main(void) {
    if (signal(SIGUSR1, sig_usr1) == SIG_ERR) printf("signal(SIGUSR1)
error");
    if (signal(SIGALRM, sig_alm) == SIG_ERR) printf("signal(SIGALRM)
error");
    pr_mask("starting main: ");
    if (sigsetjmp(jmpbuf, 1)) {
        pr_mask("ending main: ");
        exit(0);
    }
    canjump = 1;
    for (;;) pause();
}

```

```

$ ./a.out &
[1] 8874
starting main:
$ kill -SIGUSR1 8874
starting sig_usr1: SIGUSR1
in sig_alm: SIGUSR1 SIGALRM
finishing sig_usr1: SIGUSR1
ending main:  # SIGUSR1 if savesigs is 0 (i.e. mask in the sig_usr1
function)

```

## The System Function

- In [Ch06](#)
  - a simple system implementation is introduced
  - but it does not handle signals
- POSIX.1 requires the system function
  - ignore [SIGINT/SIGQUIT](#)
    - these signals should be sent only to child process
  - block [SIGCHLD](#)
    - parent will not confuse between the termination of child and other child
- System Implementation
  - preserve original signal action for [SIGINT](#) and [SIGQUIT](#)
  - block [SIGCHLD](#)
  - child
    - restore handler
    - unblock [SIGCHLD](#)
    - execute command
  - parent
    - wait for child
  - restore handler and unblock [SIGCHLD](#)

```

int system(const char *cmdstring) {
    pid_t pid;
    int status;
    struct sigaction ignore, saveintr, savequit;
    sigset_t chldmask, savemask;
    if (cmdstring == NULL) return 1;
    ignore.sa_handler = SIG_IGN;
    sigemptyset(&ignore.sa_mask);
    ignore.sa_flags = 0;
    if (sigaction(SIGINT, &ignore, &saveintr) < 0) return -1;
    if (sigaction(SIGQUIT, &ignore, &savequit) < 0) return -1;
    sigemptyset(&chldmask);
    sigaddset(&chldmask, SIGCHLD);
    if (sigprocmask(SIG_BLOCK, &chldmask, &savemask) < 0) return -1;
    if ((pid = fork()) < 0) {
        status = -1;
    } else if (pid == 0) {
        sigaction(SIGINT, &saveintr, NULL);
        sigaction(SIGQUIT, &savequit, NULL);
        sigprocmask(SIG_SETMASK, &savemask, NULL);
        execl("/bin/sh", "sh", "-c", cmdstring, (char *)0);
        _exit(127);
    } else {
        while (waitpid(pid, &status, 0) < 0) {
            if (errno != EINTR) {
                status = -1;
                break;
            }
        }
    }
    if (sigaction(SIGINT, &saveintr, NULL) < 0) return -1;
    if (sigaction(SIGQUIT, &savequit, NULL) < 0) return -1;
    if (sigprocmask(SIG_SETMASK, &savemask, NULL) < 0) return -1;
    return status;
}

```

- Job Control Signals

- except **SIGCHLD**, most applications don't handle these signals
- vi editor will save/restore terminal state when the process is stopped/continued

Signal	Description
<b>SIGCHLD</b>	Child stopped or terminated
<b>SIGCONT</b>	Continue if stopped
<b>SIGSTOP</b>	Stop signal
<b>SIGTSTP</b>	Stop typed at terminal
<b>SIGTTIN</b>	Terminal input for background process

Signal	Description
<b>SIGTTOU</b>	Terminal output for background process