

Threads

Table of Contents

- [Threads](#)
 - [Table of Contents](#)
 - [Overview](#)
 - [Thread Creation](#)
 - [Thread Termination](#)
 - [Thread Synchronization](#)
 - [Deadlock Avoidance](#)
 - [Reader-Writer Lock](#)
 - [Condition Variable](#)
 - [Barrier](#)

Overview

- Introduction
 - multiple tasks within a single process
 - they can access to the same process components
 - synchronization for shared resources
 - A Thread Consists of
 - thread ID
 - register values
 - stack content
 - a signal mask
 - an errno variable
 - scheduling priority and policy
 - thread specific data
 - Unix Thread Standard
 - POSIX.1-2001, known as pthreads
 - Linux Implementation of POSIX Threads
 - via clone system call
 - two flavors
 - LinuxThreads
 - Native POSIX Thread Library (NPTL)
 - better conformance to POSIX.1
 - e.g. POSIX.1 requires threads of a process obtaining same PID by `getpid()`, but LinuxThreads does not follow it
 - better performance
 - require supports from the C library and the kernel
 - both are 1:1 thread model
 - i.e. each thread maps to a kernel scheduling entity
- Thread Identification
 - `pthread_t` data type, and it can be a structure

- test equivalence
 - `int pthread_equal(pthread_t tid1, pthread_t tid2);`
 - return: nonzero if equal, otherwise zero
- get the current thread ID
 - `pthread_t pthread_self(void);`

Thread Creation

```
int pthread_create(pthread_t *thread,
                  const pthread_attr_t *attr,
                  void *(*start_routine) (void *),
                  void *arg);
```

- Create Function
 - `thread`: be declared previously
 - `attr`: customize thread attributes
 - `start_routine`: routine function
 - `arg`: arguments passed to `start_routine`
 - return: 0 if OK, `errno` on failure
- Note
 - no guarantee which thread runs first
 - threads have access to the process address space
 - threads inherit followings from the calling process
 - floating-point environment
 - signal mask
 - return error code, not use `errno` variable
 - per thread copy of `errno` is still provided for compatibility
- Example
 - code

```
pthread_t ntid;
void printids(const char *s) {
    pid_t pid = getpid();
    pthread_t tid = pthread_self();
    printf("%s pid %u tid %lu (%#lx)\n", s, pid, tid, tid);
}

void *thr_fn(void *arg) {
    printids("new thread: ");
    return NULL;
}

int main(void) {
```

```

int err = pthread_create(&tid, NULL, thr_fn, NULL);
if (err != 0) {
    return 1;
}
printids("main thread:");
sleep(1);
return 0;
}

```

- result
 - differs on different platforms
 - `pthread_t` may be not an integer
 - `getpid()` may return different values

```

main thread: pid 17242 tid 873604928 (0x34122740)
new thread:  pid 17242 tid 865298176 (0x33936700)

```

Thread Termination

- Terminate
 - entire process
 - `exit(3)`, `_Exit(2)` or `_exit(2)` was called
 - received signal with default action of terminating process
 - signal thread
 - `pthread_exit(3)` was called
 - return from the start routine
 - canceled by another thread in the same process
- Termination Status
 - `wait(2)` can retrieve exit status of a thread
 - `pthread_join(3)`: `int pthread_join(pthread_t thread, void **retval);`
 - return: 0 OK, `errno` on failure
 - suspends the calling thread unless the target has already terminated
 - `retval` stores exit status if it is not NULL
 - the target thread is then placed in a detached state
 - storage of a thread can be released right on its termination
 - `pthread_detach(3)`: `int pthread_detach(pthread_t thread);`
 - return: 0 OK, `errno` on failure
 - cannot be joined, such attempts will return `EINVL`
 - example

```

void errquit(const char *msg) {
    perror(msg);
    exit(-1);
}

void *thr_fn1(void *arg) {
    printf("thread 1 returning\n");
    return ((void *)1);
}

void *thr_fn2(void *arg) {
    printf("thread 2 exiting\n");
    pthread_exit((void *)2);
}

int main(void) {
    int err;
    pthread_t tid1, tid2;
    int ret;

    err = pthread_create(&tid1, NULL, thr_fn1, NULL);
    if (err != 0) errquit("create thread 1");
    err = pthread_create(&tid2, NULL, thr_fn2, NULL);
    if (err != 0) errquit("create thread 2");
    err = pthread_join(tid1, (void **)&ret);
    if (err != 0) errquit("join thread 1");
    printf("thread 1 exit code %d\n", ret);
    err = pthread_join(tid2, (void **)&ret);
    if (err != 0) errquit("join thread 2");
    printf("thread 2 exit code %d\n", ret);
}

```

- `void *`
 - can be used to pass a data structure
 - however, the data structure should not be placed on the stack
 - it might be reused by other threads when the thread is terminated
- Canceling a Thread
 - `pthread_cancel(3)`: `int pthread_cancel(pthread_t thread);`
 - return: 0 OK, `errno` on failure
 - sends a cancellation request to `thread`
 - similar to `thread` calls `pthread_exit(PTHREAD_CANCELED)`
 - `thread` can select to ignore or control how it is canceled
 - **does not wait for the thread to terminate**
- Cleanup Functions
 - recall: `atexit(3)` can register functions executed when termination
 - `pthread_cleanup_push(3)`: `void pthread_cleanup_push(void (*routine)(void *), void *arg);`
 - `pthread_cleanup_pop(3)`: `void pthread_cleanup_pop(int execute);`

- registered routines is executed when
 - calls `pthread_exit(3)`
 - responding to a cancellation request
 - calls `pthread_cleanup_pop(3)` with a nonzero argument
 - if zero argument, it just remove the routine on stack top
- Comparison of Process and Thread Primitives

Process	Thread Primitive	Description
<code>fork(2)</code>	<code>pthread_create(3)</code>	Create a new flow of control
<code>exit(3)</code>	<code>pthread_exit(3)</code>	Exit from an existing flow of control
<code>waitpid(2)</code>	<code>pthread_join(3)</code>	Get exit status from flow of control
<code>atexit(3)</code>	<code>pthread_cleanup_push(3)</code>	Register function to be called at exit from flow of control
<code>getpid(2)</code>	<code>pthread_self(3)</code>	Get ID for flow of control
<code>abort(3)</code>	<code>pthread_cancel(3)</code>	Request abnormal termination of flow of control

Thread Synchronization

- Mutexes
 - `pthread_mutex_init(3p):int pthread_mutex_init(pthread_mutex_t *restrict mutex, const pthread_mutexattr_t *restrict attr);`
 - alternative `pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER`
 - `pthread_mutex_destroy(3p):int pthread_mutex_destroy(pthread_mutex_t *mutex);`
 - `pthread_mutex_trylock(3p):int pthread_mutex_trylock(pthread_mutex_t *mutex);`
 - `pthread_mutex_lock(3p):int pthread_mutex_lock(pthread_mutex_t *mutex);`
 - `pthread_mutex_unlock(3p):int pthread_mutex_unlock(pthread_mutex_t *mutex);`
 - return: 0 OK, error number on failure
 - example: protect data structure

```
// foo_alloc allocate memory and initialize f_lock

void foo_hold(struct foo *fp) {
    pthread_mutex_lock(&fp->f_lock);
    fp->f_count++;
    pthread_mutex_unlock(&fp->f_lock);
}

void foo_rele(struct foo *fp) {
    pthread_mutex_lock(&fp->f_lock);
    if (--fp->f_count == 0) {
        pthread_mutex_unlock(&fp->f_lock);
        pthread_mutex_destroy(&fp->f_lock);
    }
}
```

```

    free(fp);
} else {
    pthread_mutex_unlock(&fp->f_lock);
}
}

```

Deadlock Avoidance

- Solution
 - lock performs in the same order
 - `pthread_mutex_trylock(3p)` can be used to check locks

Reader-Writer Lock

- Reader-Writer Lock
 - similar to mutexes, but higher degree of parallelism
 - mutexes: locked or unlocked
 - reader-writer lock: locked in read mode, locked in write mode, or unlocked
 - **multiple reader locks can be acquired simultaneously**
 - **but only one can lock in write mode**
 - **if a reader/writer locks, then the coming writer/reader must wait until it unlocks**
 - `pthread_rwlock_init(3p): int pthread_rwlock_init(pthread_rwlock_t *restrict rwlock, const pthread_rwlockattr_t *restrict attr);`
 - alternative `pthread_rwlock_t rwlock = PTHREAD_RWLOCK_INITIALIZER;`
 - `pthread_rwlock_destroy(3p): int pthread_rwlock_destroy(pthread_rwlock_t *rwlock);`
 - `pthread_rwlock_tryrdlock(3p): int pthread_rwlock_tryrdlock(pthread_rwlock_t *rwlock);`
 - `pthread_rwlock_rdlock(3p): int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock);`
 - `pthread_rwlock_trywrlock(3p): int pthread_rwlock_trywrlock(pthread_rwlock_t *rwlock);`
 - `pthread_rwlock_wrlock(3p): int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);`
 - `pthread_rwlock_unlock(3p): int pthread_rwlock_unlock(pthread_rwlock_t *rwlock);`
 - return: 0 OK, error number on failure

Condition Variable

- Condition Variable
 - condition is protected by a mutex
 - a thread must lock the mutex to change the condition state
 - allow a thread to wait in a race-free way for arbitrary conditions to occur

- `pthread_cond_init(3p):int pthread_cond_init(pthread_cond_t *restrict cond, const pthread_condattr_t *restrict attr);`
 - alternative `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`
- `pthread_cond_destroy(3p):int pthread_cond_destroy(pthread_cond_t *cond);`
- `pthread_cond_wait(3p):int pthread_cond_wait(pthread_cond_t *restrict cond, pthread_mutex_t *restrict mutex);`
 1. unlock `mutex` (`mutex` must be locked when it is called)
 2. wait for `cond` to occur (thread start to sleep)
 3. if being signaled, lock `mutex`
- `pthread_cond_timedwait(3p):int pthread_cond_timedwait(pthread_cond_t *restrict cond, pthread_mutex_t *restrict mutex, const struct timespec *restrict abstime);`
 - occur error `ETIMEDOUT` if `abstime` is passed
 - absolute expire time may need `gettimeofday(2)` to set

```
struct timespec {
    long int tv_sec;
    long int tv_nsec;
};
```

- `pthread_cond_broadcast(3p):int pthread_cond_broadcast(pthread_cond_t *cond);`
 - wake up all waiting threads
- `pthread_cond_signal(3p):int pthread_cond_signal(pthread_cond_t *cond);`
 - wake up one waiting thread
 - POSIX.1 allows the implementation wakes up more than one threads
 - waked up threads have to contend for the mutex
 - return: 0 OK, error number on failure

- example: producer and consumer

```
void process_msg(void) {
    struct msg *mp;
    for (;;) {
        pthread_mutex_lock(&qlock);
        /* spurious wakeup */
        /* because other threads may get the resource and set workq to
        NULL */
        while (workq == NULL) pthread_cond_wait(&qready, &qlock);
        mp = workq;
```

```

        workq = mp->m_next;
        pthread_mutex_unlock(&qlock);
        /* now process the message mp */
    }
}

void enqueue_msg(struct msg *mp) {
    pthread_mutex_lock(&qlock);
    mp->m_next = workq;
    workq = mp;
    pthread_mutex_unlock(&qlock);
    pthread_cond_signal(&qready);
}

```

```

process: (work, NULL) and wait
enqueue: (work, 1)
enqueue: (work, 2) -> (1)
enqueue: (work, 3) -> (2) -> (1)
process: (work, 2) -> (1)
enqueue: (work, 4) -> (2) -> (1)
process: (work, 2) -> (1)
process: (work, 1)
process: (work, NULL) and wait

```

- example: job queue

```

#define N_WORKERS 3
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
std::list<Job> jobqueue;

int do_the_job(long id, int ch) {
    if (ch == -1) return -1;
    printf("worker-%ld: %c\n", id, ch);
    return 0;
}

void *worker_main(void *arg) {
    long id = (long)arg;
    printf("# worker-%ld created\n", id);
    while (1) {
        Job j;
        pthread_mutex_lock(&mutex);
        pthread_cond_wait(&cond, &mutex);
        // signaled, means at least one job
        j = jobqueue.front();
        if (j.getId() == 0 || pthread_equal(pthread_self(), j.getId())) {
            jobqueue.pop_front();
        }
#ifdef ORDERED // follow string order

```



```

        if (do_the_job(id, j.getChar()) < 0) {
            pthread_mutex_unlock(&mutex);
            break;
        }
    #endif
    } else {
        pthread_mutex_unlock(&mutex);
        continue;
    }
    pthread_mutex_unlock(&mutex);
    #ifndef ORDERED // may not follow string order
        if (do_the_job(id, j.getChar()) < 0) break;
    #endif
    }
    printf("# worker-%ld terminated\n", id);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t workers[N_WORKERS];
    if (argc < 2) {
        fprintf(stderr, "usage: %s input-string\n", argv[0]);
        return -1;
    }
    // create workers
    for (int i = 0; i < N_WORKERS; i++) {
        if (pthread_create(&workers[i], NULL, worker_main, (void *)
(long)i) != 0) {
            fprintf(stderr, "create worker[%d] failed\n", i);
            exit(-1);
        }
    }
    // create jobs, signal every time a new job is pushed in queue
    for (char *ptr = argv[1]; *ptr; ptr++) {
    #ifndef ASSIGNID
        Job j(*ptr, workers[(ptr - argv[1]) % N_WORKERS]);
    #else
        Job j(*ptr);
    #endif
        pthread_mutex_lock(&mutex);
        jobqueue.push_back(j);
        pthread_mutex_unlock(&mutex);
        pthread_cond_signal(&cond);
    }
    // terminate workers
    for (int i = 0; i < N_WORKERS; i++) {
    #ifndef ASSIGNID
        Job j(-1, workers[i]);
    #else
        Job j(-1);
    #endif
        pthread_mutex_lock(&mutex);
        jobqueue.push_back(j);
        pthread_mutex_unlock(&mutex);

```

```

    pthread_cond_signal(&cond);
}
// process all jobs
size_t jobs;
do {
    pthread_mutex_lock(&mutex);
    jobs = jobqueue.size();
    pthread_mutex_unlock(&mutex);
    pthread_cond_signal(&cond);
} while (jobs > 0);
// wait for all workers
for (int i = 0; i < N_WORKERS; i++) {
    void *ret;
    pthread_join(workers[i], &ret);
}
return 0;
}

```

Barrier

- `pthread_barrier_init(3p)`: `int pthread_barrier_init(pthread_barrier_t *restrict barrier, const pthread_barrierattr_t *restrict attr, unsigned count);`
- `pthread_barrier_destroy(3p)`: `int pthread_barrier_destroy(pthread_barrier_t *barrier);`
- `pthread_barrier_wait(3p)`: `int pthread_barrier_wait(pthread_barrier_t *barrier);`
 - return: 0 OK, error number on failure
- Example

```

#define N 5
#ifdef HAS_BARRIER
static pthread_barrier_t barrier;
#endif

void *worker(void *arg) {
    long i, id = (long)arg;
    for (i = 0; i < id + 1; i++) {
        fprintf(stderr, "%ld", id + 1);
    }
    fprintf(stderr, "[%ld/done]\n", id + 1);
#ifdef HAS_BARRIER
    pthread_barrier_wait(&barrier);
#endif
    return NULL;
}

```

```
int main() {
    long i;
    pthread_t tid;
#ifdef HAS_BARRIER
    pthread_barrier_init(&barrier, NULL, N + 1);
#endif
    for (i = 0; i < N; i++) {
        if (pthread_create(&tid, NULL, worker, (void *)i) != 0) {
            fprintf(stderr, "pthread_create failed.\n");
            return -1;
        }
    }
#ifdef HAS_BARRIER
    pthread_barrier_wait(&barrier);
    pthread_barrier_destroy(&barrier);
#endif
    fprintf(stderr, "all done.\n");
    return 0;
}
```