

Daemon Process

- [Daemon Process](#)
 - [Introduction](#)
 - [Daemon Conventions](#)
 - [Client Server Model](#)
 - [Coding Rules](#)
 - [Error Logging](#)
 - [Single Instance Daemon](#)

Introduction

- Processes Live for a Long Time
 - often started when the system is up and terminate when it is down
 - do not have a controlling terminal
 - run in background
 - partial output `ps ajx`
 - `x`: process without terminal
 - `j`: show job related information
 - processes enclosed by `[]` are kernel processes

Daemon Conventions

- Lock File
 - if daemon uses a lock file, it is usually stored in `/var/run/name.pid`
 - use to check the existence of a running daemon
 - might need superuser permissions to create a file here
- Configuration Options
 - usually stored in `/etc`, named `name.conf`
 - daemon usually reads it only if it starts
 - some daemon will read configuration again when received `SIGHUP`
- Start
 - start from `/etc/rc*` or `/etc/init.d/*`

Client Server Model

- daemon is common to use as server
- e.g.
 - client: `syslog` library call
 - server: `syslogd` daemon

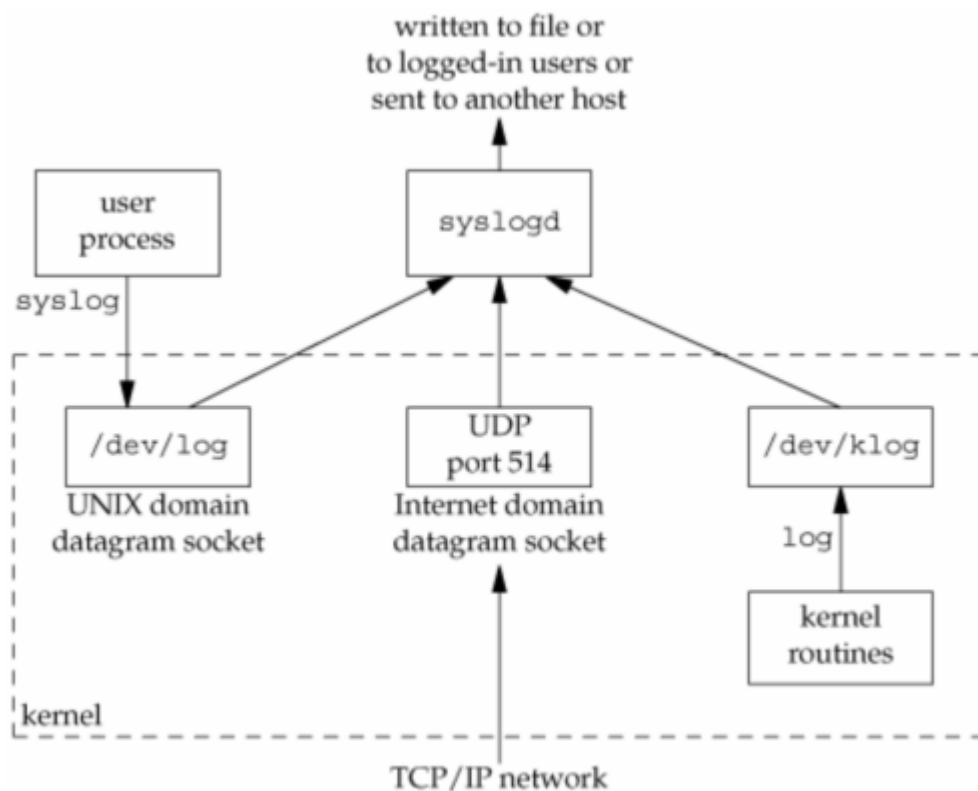
Coding Rules

1. `umask(2)` to reset permission masks
2. `fork(2)` and parent exit
 - not block the shell

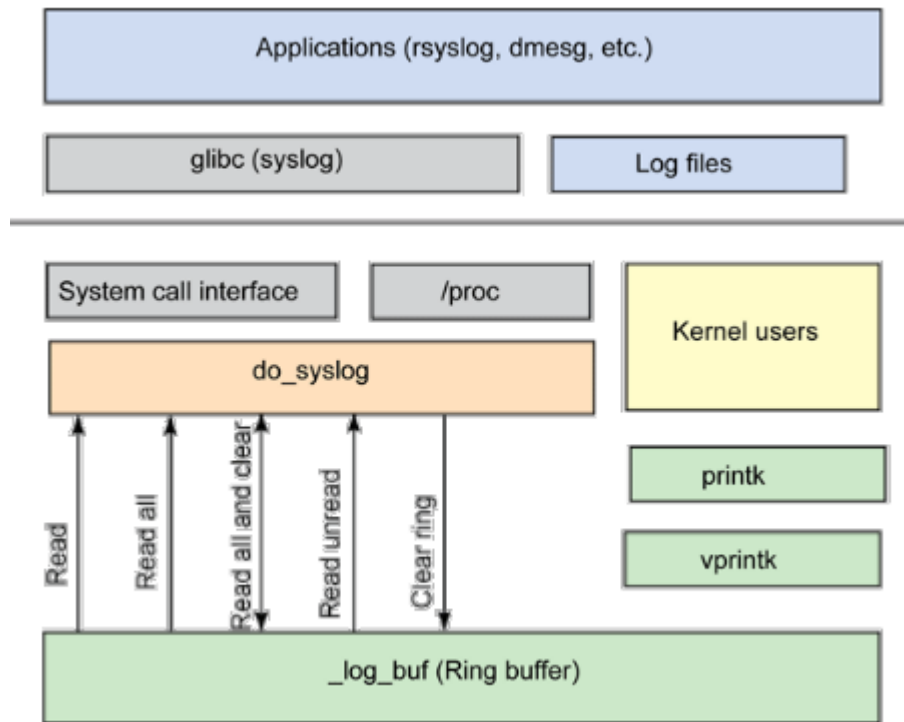
- not a process group leader
 - child inherits the process group ID of the parent
 - but gets a new process ID
 - can be a session leader
3. `setsid(2)`
 - become a session leader
 - become a process group leader
 - no controlling terminal
 - System V systems suggest to `fork(2)` again and terminate the parent
 - child continues as the daemon, and not a session leader
 - prevent the child from acquiring controlling terminal
 4. change the current working directory, may be root directory
 5. close unused file descriptors, and redirect {0, 1, 2} to `/dev/null`
 6. setup log files or log systems

Error Logging

- Logging
 - daemon cannot write error message to standard error
 - a centralized place to collect logs from all daemons
 - `syslog` facility initially proposed in the BSD system
 - many systems derived the design from `syslogd`
 - e.g. Linux's `rsyslogd`: a reliable and extended of `syslogd`



- Linux Kernel Logging Stack
 - `syslog(2)`: read and clear the kernel log buffer
 - `syslog(3)`: write log to log system



- `syslog(3)`

- `void openlog(const char *ident, int option, int facility);`
- `void syslog(int priority, const char *format, ...);`
- `void closelog(void);`
- the call to `openlog(3)` is optional
 - it will be called on the first call to `syslog(3)`
- the call to `closelog(3)` is optional
 - it simply close the descriptor
- `priority` argument can be specified by facility and level together
- option

Option	Description
<code>LOG_CONS</code>	Write to the console if there is an error while sending to the system logger
<code>LOG_NDELAY</code>	Open the connection immediately instead of opening when the first message is logged
<code>LOG_NOWAIT</code>	Do not wait for child processes that may have been created while logging the message
<code>LOG_ODELAY</code>	Delay the connection until <code>syslog(3)</code> is called
<code>LOG_PERROR</code>	(Not in POSIX.1-2001 or POSIX.1-2008.) Also log the message to <code>stderr</code>
<code>LOG_PID</code>	Include the caller's PID with each message

- facility

Facility	Description
<code>LOG_AUTH</code>	security/authorization messages
<code>LOG_AUTHPRIV</code>	security/authorization messages (private)

Facility	Description
LOG_CRON	clock daemon (cron and at)
LOG_DAEMON	system daemons without separate facility value
LOG_FTP	ftp daemon
LOG_KERN	kernel messages (these can't be generated from user processes)
LOG_LOCAL0~7	reserved for local use
LOG_LPR	line printer subsystem
LOG_MAIL	mail subsystem
LOG_NEWS	USENET news subsystem
LOG_SYSLOG	messages generated internally by <code>syslogd(8)</code>
LOG_USER	generic user-level messages(default)
LOG_UUCP	UUCP subsystem

- level

Level	Description
LOG_EMERG	system is unusable
LOG_ALERT	action must be taken immediately
LOG_CRIT	critical conditions
LOG_ERR	error conditions
LOG_WARNING	warning conditions
LOG_NOTICE	normal, but significant, condition
LOG_INFO	informational message
LOG_DEBUG	debug-level message

- `setlogmask(3): int setlogmask(int mask);`
 - can be used to restrict logging to specified levels only.
- Example

```
// line printer spooler daemon
openlog("lpd", LOG_PID, LOG_LPR);
syslog(LOG_ERR, "open error for %s: %m", filename);
// an almost equivalent implementation
syslog(LOG_ERR | LOG_LPR, "open error for %s: %m", filename);
```

Single Instance Daemon

- Single Instance
 - some daemons are implemented so that only a single copy is allowed
 - e.g. the daemon might need exclusive access to a device
 - the cron daemon
- Codes to Check Daemon Existence

```
int already_running(d) {
    int fd;
    char buf[16];
    if ((fd = open(LOCKFILE, O_RDWR | O_CREAT, LOCKMODE)) < 0) {
        syslog(LOG_ERR, "can't open %s: %s", LOCKFILE, strerror(errno));
        exit(1);
    }
    if (lockfile(fd) < 0) {
        if (errno == EACCES || errno == EAGAIN) {
            close(fd);
            return 1;
        }
        syslog(LOG_ERR, "can't lock %s: %s", LOCKFILE, strerror(errno));
        exit(1);
    }
    ftruncate(fd, 0);
    sprintf(buf, "%ld\n", (long)getpid());
    write(fd, buf, strlen(buf));
    return 0;
}
```